

---

# DTS207TC Database Development and Design

## Lecture 12

### Chap 31. Information Retrieval

Di Zhang, Autumn 2025

# Outline

---

- History
- IF-IDF
- Pagerank
- Revert-Index
- Web Crawler

# The History of Search Engines

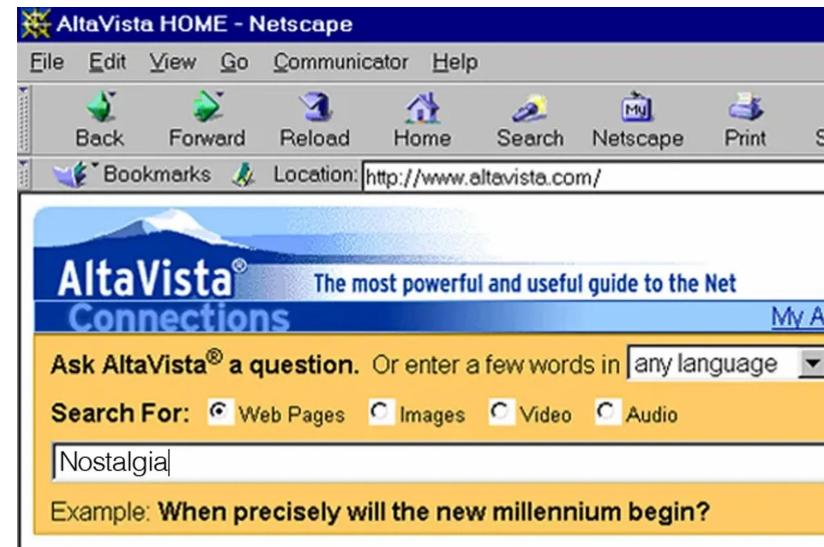
- The Early Days: Pre-Web & The First Steps
- Before Google: The Pioneers (1990-1995)
  - The Problem: The early internet was a collection of disconnected files. No easy way to find anything.
  - Archie (1990): The first "search engine." It simply indexed the names of files on public FTP servers.
  - Gopher & Veronica/Jughead: Searched menu-based systems, not full text.
  - The Web Explosion: The invention of the World Wide Web (WWW) created a urgent need for better navigation tools.
  - Wandex (1993): The first web robot and indexer, cataloging the early web.



# The Rise of Web Crawlers & The First Giants

The Crawler Revolution (1994-1998)

- WebCrawler (1994): The first full-text search engine that allowed users to search for any word on any webpage.
- The Directory Approach - Yahoo!: Founded by humans who manually categorized websites. It was a curated web directory, not a true crawler-based engine.
- The Powerhouses: AltaVista & Excite:
  - AltaVista: Launched by DEC, it was fast, powerful, and supported natural language queries and advanced search operators. It was the king of search for years.
  - The Common Flaw: These early engines relied heavily on basic keyword matching, making them easy to spam with irrelevant pages.



# The Google Revolution: PageRank Changes Everything

The Game Changer: Google (1998-Present)

- The Breakthrough: PageRank Algorithm.
  - Founded by Larry Page and Sergey Brin at Stanford.
  - Core Idea: A webpage is important if other important pages link to it. It treated links as "votes."
- Why It Won:
  - Relevance: Delivered vastly more relevant and higher-quality results.
  - Speed: Incredibly fast.
  - Simplicity: The clean, uncluttered homepage was a stark contrast to the "portal" look of competitors.
- The Domination: Google's superior technology quickly made it the default search engine for the world.



# The Modern Era & The AI-Powered Future

## Beyond Keywords: Voice, Personalization, and AI

- Mobile & Voice Search: Search moved from the desktop to everywhere. Queries became conversational ("OK Google, where's the nearest coffee shop?").
- Personalization: Search results are tailored to your location, search history, and preferences.
- Featured Snippets & "Answer Engines": Google started pulling information directly to the top of the page, aiming to answer your question without a click.
- The AI & LLM Revolution (Now & Future):
  - BERT & MUM: Google's AI models understand context and nuance in queries like never before.
  - Generative AI & ChatGPT: The rise of Large Language Models (LLMs) is shifting search from a "list of links" to a conversational experience. The future is about getting synthesized, direct answers and engaging in a dialogue.



# Information Retrieval Systems

- **Information retrieval (IR)** systems use a simpler data model than database systems
  - Information organized as a collection of documents
  - Documents are unstructured, no schema
- Information retrieval locates relevant documents, on the basis of user input such as keywords or example documents
  - e.g., find documents containing the words “database systems”
- Can be used even on textual descriptions provided with non-textual data such as images
- Web search engines are the most familiar example of IR systems

- Differences from database systems
  - IR systems don't deal with transactional updates (including concurrency control and recovery)
  - Database systems deal with structured data, with schemas that define the data organization
  - IR systems deal with some querying issues not generally addressed by database systems
    - Approximate searching by keywords
    - Ranking of retrieved answers by estimated degree of relevance

# Keyword Search

- In **full text** retrieval, all the words in each document are considered to be keywords.
  - We use the word **term** to refer to the words in a document
- Information-retrieval systems typically allow query expressions formed using keywords and the logical connectives *and*, *or*, and *not*
  - *Ands* are implicit, even if not explicitly specified
- Ranking of documents on the basis of estimated relevance to a query is critical
  - Relevance ranking is based on factors such as
    - Term frequency
      - Frequency of occurrence of query keyword in document
    - Inverse document frequency
      - How many documents the query keyword occurs in
        - Fewer → give more importance to keyword
    - Hyperlinks to documents
      - More links to a document → document is more important

# Relevance Ranking Using Terms

- **TF-IDF** (Term frequency/Inverse Document frequency)  
ranking:

$$TF(t, d) = \frac{\text{number of times } t \text{ appears in } d}{\text{total number of terms in } d}$$

$$IDF(t) = \log \frac{N}{1 + df}$$

$$TF - IDF(t, d) = TF(t, d) * IDF(t)$$

$tf(t, d)$

	blue	bright	can	see	shining	sky	sun	today
1	1/2	0	0	0	0	1/2	0	0
2	0	1/3	0	0	0	0	1/3	1/3
3	0	1/3	0	0	0	1/3	1/3	0
4	0	1/6	1/6	1/6	1/6	0	1/3	0

$idf(t, D)$

	blue	bright	can	see	shining	sky	sun	today
	0.602	0.125	0.602	0.602	0.602	0.301	0.125	0.602

×

$$tfidf(t, d, D) = tf(t, d) \cdot idf(t, D)$$

	blue	bright	can	see	shining	sky	sun	today
1	0.301	0	0	0	0	0.151	0	0
2	0	0.0417	0	0	0	0	0.0417	0.201
3	0	0.0417	0	0	0	0.100	0.0417	0
4	0	0.0209	0.100	0.100	0.100	0	0.0417	0

- TF-IDF: Multiply TF and IDF scores, use to rank importance of words within documents
- Most important word for each document is highlighted

→

# TF-IDF is the measure of importance in information theory\*

- The Goal:

- Find term weights that make each document stand out from the corpus average.

- The Math:

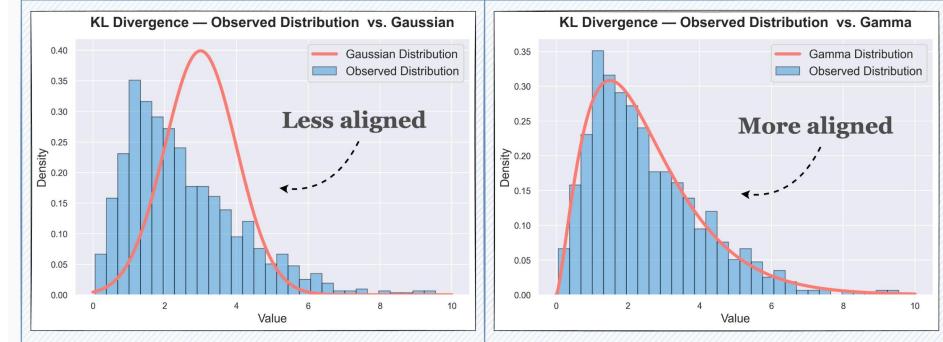
- We model this by maximizing the difference between:
  - A term's local distribution in a doc,  $P(t|d)$  ( $\approx$  TF)
  - Its global background distribution,  $P(t)$

- The Insight:

- The optimal measure for this difference is KL Divergence:

$$D_{KL}(P_d||P_C) = \sum_w P_d(w) \log \frac{P_d(w)}{P_C(w)}$$

## KL Divergence



KL Divergence = 78.9



KL Divergence = 46.3



$$\text{Score}(w) \propto \underbrace{P_d(w)}_{\text{TF}} \times \underbrace{\log \frac{N}{df(w)}}_{\text{IDF}}$$

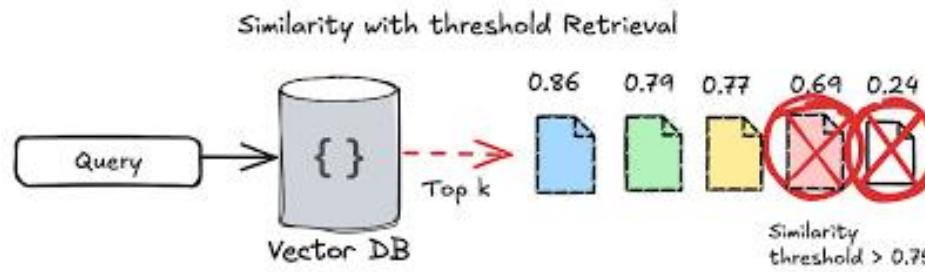
# Relevance Ranking Using Terms (Cont.)

---

- Most systems add to the above model
  - Words that occur in title, author list, section headings, etc. are given greater importance
  - Words whose first occurrence is late in the document are given lower importance
  - Very common words such as “a”, “an”, “the”, “it” etc. are eliminated
    - Called **stop words**
  - **Proximity:** if keywords in query occur close together in the document, the document has higher importance than if they occur far apart
- Documents are returned in decreasing order of relevance score
  - Usually only top few documents are returned, not all

# Similarity Based Retrieval

- Similarity based retrieval - retrieve documents similar to a given document
  - Similarity may be defined on the basis of common words
    - E.g., find  $k$  terms in A with highest  $TF(d, t) / n(t)$  and use these terms to find relevance of other documents.
- **Relevance feedback:** Similarity can be used to refine answer set to keyword query
  - User selects a few relevant documents from those retrieved by keyword query, and system finds other documents similar to these
- Vector space model: define an  $n$ -dimensional space, where  $n$  is the number of words in the document set.
  - Vector for document  $d$  goes from origin to a point whose  $i^{\text{th}}$  coordinate is  $TF(d, t) / n(t)$
  - The cosine of the angle between the vectors of two documents is used as a measure of their similarity.



# Relevance Using Hyperlinks

---

- Number of documents relevant to a query can be enormous if only term frequencies are taken into account
- Using term frequencies makes “spamming” easy
  - E.g., a travel agency can add many occurrences of the words “travel” to its page to make its rank very high
- Most of the time people are looking for pages from popular sites
- Idea: use popularity of Web site (e.g., how many people visit it) to rank site pages that match given keywords
- Problem: hard to find actual popularity of site
  - Solution: next slide

# Relevance Using Hyperlinks (Cont.)

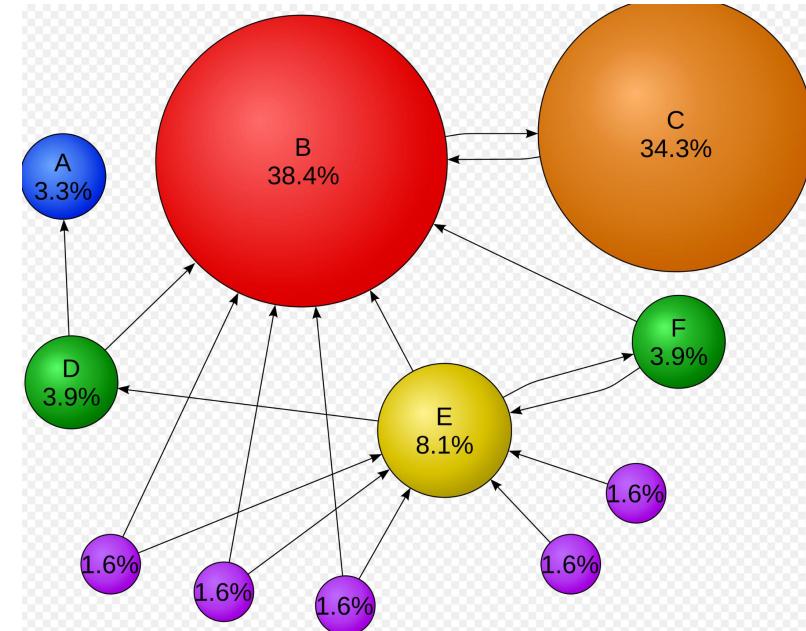
- Solution: use number of hyperlinks to a site as a measure of the popularity or **prestige** of the site
  - Count only one hyperlink from each site (why? - see previous slide)
  - Popularity measure is for site, not for individual page
    - But, most hyperlinks are to root of site
    - Also, concept of “site” difficult to define since a URL prefix like cs.yale.edu contains many unrelated pages of varying popularity
- Refinements
  - When computing prestige based on links to a site, give more weight to links from sites that themselves have higher prestige
    - Definition is circular
    - Set up and solve system of simultaneous linear equations
  - Above idea is basis of the Google **PageRank** ranking mechanism

# Relevance Using Hyperlinks (Cont.)

- Connections to **social networking** theories that ranked prestige of people
  - E.g., the president of the U.S.A has a high prestige since many people know him
  - Someone known by multiple prestigious people has high prestige
- Hub and authority based ranking
  - A **hub** is a page that stores links to many pages (on a topic)
  - An **authority** is a page that contains actual information on a topic
  - Each page gets a **hub prestige** based on prestige of authorities that it points to
  - Each page gets an **authority prestige** based on prestige of hubs that point to it
  - Again, prestige definitions are cyclic, and can be got by solving linear equations
  - Use authority prestige when ranking answers to a query

# Introduction to PageRank

- What is PageRank?
  - A link analysis algorithm developed by Google founders Larry Page and Sergey Brin at Stanford.
  - Core Idea: A web page is important if it is linked to by other important pages.
  - It assigns a numerical weight to each element of a hyperlinked set of documents (like the World Wide Web) to measure its relative importance.
  - Originally the foundation of Google's web search ranking.
- The Intuition: The "Random Surfer" Model
  - Imagine a person randomly clicking on links.
  - The PageRank value represents the probability that this random surfer will land on a given page.
  - A page has a high rank if:
    - Many pages link to it.
    - Important pages link to it.



# How PageRank Works - The Math

- The Basic Formula:

The simplified PageRank for a page  $A$  is:

$$PR(A) = (1 - d) + d \times \left( \frac{PR(T_1)}{C(T_1)} + \dots + \frac{PR(T_n)}{C(T_n)} \right)$$

- $PR(T_n)$ : PageRank of a linking page  $T_n$ .
- $C(T_n)$ : Number of outbound links on page  $T_n$ .
- $d$  (Damping Factor): The probability that the random surfer will click a link (usually set to  $\sim 0.85$ ).
- $(1 - d)$ : The probability the surfer "gets bored" and jumps to a random page.

- It's a "Chicken and Egg" Problem:

- To know A's rank, you need to know the rank of all pages linking to A.
- This creates a system of equations that must be solved simultaneously.

# The Power Iteration Connection

- The entire web's PageRank can be represented as an eigenvector problem.

$$\mathbf{PR} = d\mathbf{M} \cdot \mathbf{PR} + \frac{(1-d)}{N} \mathbf{1}$$

- $\mathbf{PR}$ : The PageRank vector (what we want to find).
- $\mathbf{M}$ : The stochastic adjacency matrix of the web graph.
- $\mathbf{1}$ : A vector of ones.
- $N$ : Total number of pages.

## Power Iteration is the Solution

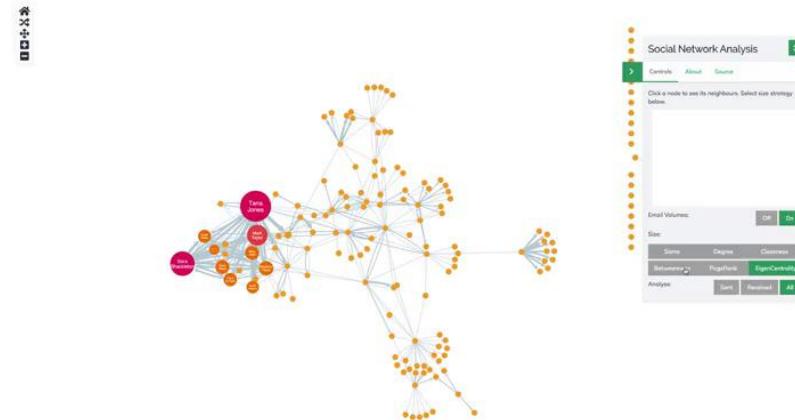
1. **Initialize:** Start with a guess (e.g.,  $PR = 1/N$  for all pages).
2. **Iterate:** Apply the formula:  $\mathbf{PR}_{k+1} = d\mathbf{M} \cdot \mathbf{PR}_k + \frac{(1-d)}{N} \mathbf{1}$
3. **Converge:** Repeat step 2 until the values stop changing significantly ( $\|\mathbf{PR}_{k+1} - \mathbf{PR}_k\| < \epsilon$ ).

- Why it Works:

- Power Iteration is a numerical method for finding the dominant eigenvector of a matrix.
- The PageRank vector is the dominant eigenvector of the modified web matrix.
- This makes the computation scalable to billions of web pages.

# Beyond Web Search - Applications

- Social Network Analysis
  - Identifying Influencers: PageRank can identify key influencers in a social graph. A user is "important" if they are followed by other important users.
  - Recommendation Systems: It can rank users or content. For example, on Twitter/X, a user's "Who to Follow" suggestions can be based on a PageRank-like score.
  - Fraud Detection: In a network of financial transactions, PageRank can help identify accounts that are central to suspicious activity.
- Other Domain Applications
  - Bioinformatics: Ranking proteins in a protein-protein interaction network to find the most critical ones.
  - Neuroscience: Identifying key regions in a brain connectivity network.
  - Citation Analysis: Determining the importance of academic papers based on which other papers cite them (the original "Impact Factor" idea).
- Any System that can be Modeled as a Graph: Where the importance of a node depends on the importance of its neighbors.



# The Inverted Index: Powering Modern Search

---

- What is the Problem?
  - We have a massive collection of documents (web pages, articles, books).
  - A user queries for a specific word or phrase.
  - How do we quickly find all documents containing that term without scanning every single document?
- The Answer: The Inverted Index
  - A data structure that maps words to the list of documents they appear in.
  - Think of it as the "index" at the back of a book, but for an entire corpus.
- Analogy: A Book's Index
  - You don't read the entire book to find a topic.
  - You look up the topic in the index, which points you to the relevant pages.

# Structure of an Inverted Index

- A Simple Example:
- We have three documents:
  - Doc1: "the cat sat on the mat"
  - Doc2: "the dog played with the cat"
  - Doc3: "the mat is clean"
- Forward Lookup (Inefficient):

Document	Content
Doc1	the cat sat on the mat
Doc2	the dog played with the cat
Doc3	the mat is clean

- Inverted Index (Efficient):

Term	Posting List (Document ID : Positions)
cat	Doc1: [2], Doc2: [6]
dog	Doc2: [3]
mat	Doc1: [6], Doc3: [2]
sat	Doc1: [3]
the	Doc1: [1, 5], Doc2: [1, 5], Doc3: [1]
...	...

- Key Components:

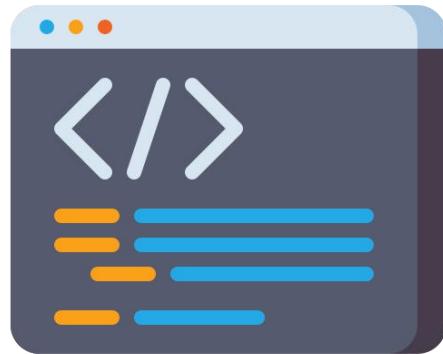
- Dictionary/Terms: The unique words (e.g., "cat", "mat").
- Posting List: The list of documents where the term appears.
- Additional Data (Optional): Term frequency, positions (for phrase queries), etc.

# Building an Index with MapReduce

- The Challenge:
  - Building an index for the entire web is impossible on a single machine. We need a distributed approach.
- Why MapReduce?
  - Automatic Parallelization: Splits the work across thousands of machines.
  - Fault Tolerance: Handles machine failures gracefully.
  - Ideal for this task: The process is inherently parallelizable.
- High-Level Process:
  - Input: A large set of raw text documents stored in a distributed file system (e.g., HDFS).
  - MapReduce Phases:
    - Map: Processes each document and emits (key, value) pairs.
    - Shuffle & Sort: Groups all values (intermediate data) by the same key.
    - Reduce: Aggregates the grouped data to form the final index.

# Demo

---

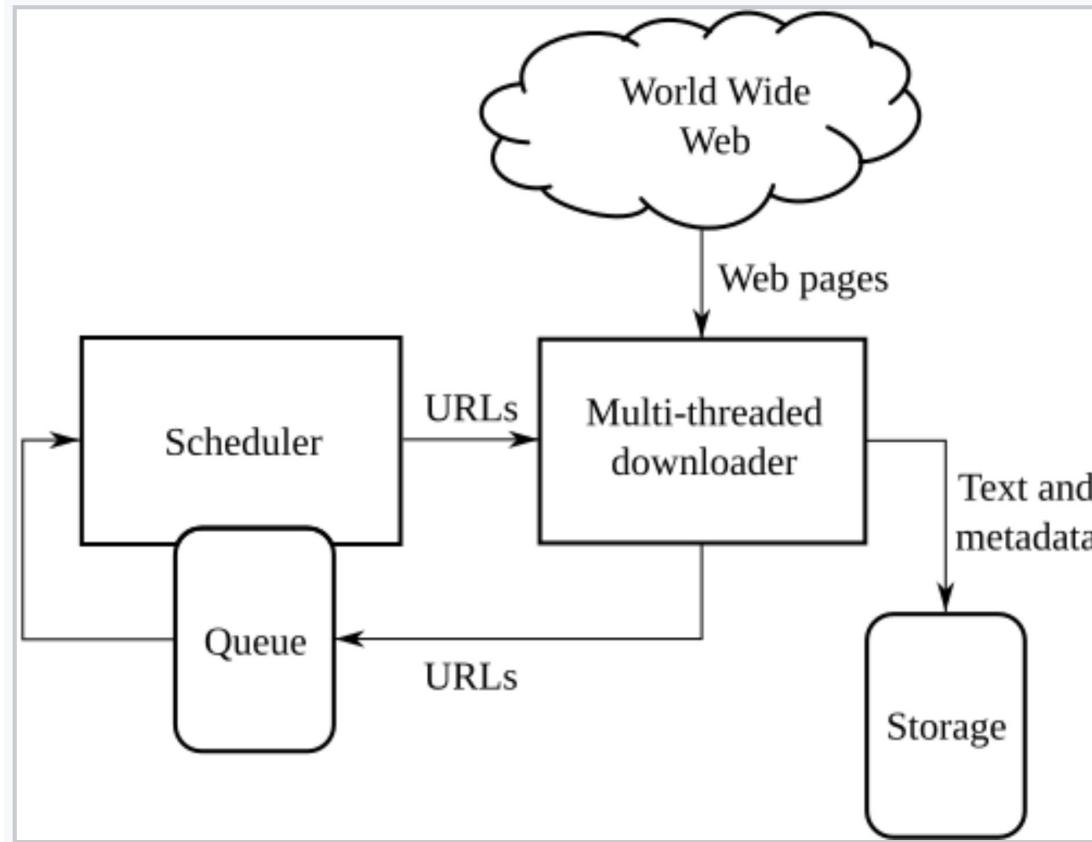


# Web Crawling

- Web crawlers systematically browse the web to collect data.
- Underlying data structures determine efficiency, scalability, and correctness.



# Architecture



# Example

---

- <https://www.octoparse.com/>
- <https://www.youtube.com/watch?v=F6CEXNb54Tl>