
DTS205TC High Performance Computing

Lecture 2 Infrastructure 2


Di Zhang, Spring 2024

History of computers



Replica of Konrad Zuse's Z3,
the first fully automatic, digital
(electromechanical) computer

1941, Germany




ENIAC was the first electronic, Turing-complete device, and performed ballistics trajectory calculations for the United States Army.

1945, US



A section of the reconstructed Manchester Baby, the first electronic stored-program computer


1948, UK



A 12-row/80-column IBM punched card from the mid-twentieth century

Von Neumann Model

Figure 1.5 The Von Neumann model



- Programs must be stored
- Sequential execution of instructions


Components

- ALU
- Memory
- Control unit
- Input / output

*Turing and Von Neumann

- A set of fixed instructions or a programming language is enough

Simplified model of x86 CPU




Demo

- https://peterhigginson.co.uk/lmc/?F5=12-Feb-24_14:04:09
- <https://justine.lol/ape.html>

Central processing unit (CPU)



Figure 5.2 Central processing unit (CPU)



Central Processing Unit (CPU)


R: Register

PC: Program Counter

IR: Instruction Register

Main memory

Figure 5.3 Main memory



Memory types


- ROM



- RAM




Structure of computer



von neumann architecture

- The main components (review of 103) :
 - CPU
 - I/O
 - Memory
 - Graphic card



Intel D945GCCR




Doom: PC Graphics Benchmark

Example: Hello world


```
1 #include <stdio.h>
2
3 int main()
4 {
5     printf("hello, world\n");
6 }
```

COMPILER	VS INTERPRETER
	
<ul style="list-style-type: none">Display All Errors after compilation, all at the same time.Faster ExecutionIntermediate object code requires memory.	<ul style="list-style-type: none">Display All Errors of each line one by one.Slower ExecutionNo intermediate code, thus less memory required.




Q: Anyone knows differences between C and Python?

Example: Hello world




1. Shell reads the hello command into memory from the keyboard.

Example: Hello world




2. OS loads the executable from disk into main memory.

Example: Hello world



3. The OS prepares the program, which then executes itself and writes the output string from memory to the display.

Speed Mismatch: Memory Wall




GFLOPS=Giga (10^9) Floating-point Operations Per Second

- Problem: In order to read and write data, the CPU must wait for the memory




An Example of Speed Mismatching


Caches matter



They can be cache at bottom!



Memory Devices in Terms of Cost & Performance



- Register: fast but expensive
- The role of cache: play the role of agent; based on the principle of locality, it can speed up the visiting
- Hierarchical cache system: Multi-Level Acceleration


Example: Summing 1-D Arrays

- Consider summing an array ‘int a[N]’

```
// Loop Over, and add every item to Addition
for(i = 0; i < Size; i++)
{
    Addition = Addition + a[i];
}
```

- It requires $N-1$ *additions*, as well as $2*(N-1)$ *reads* and $N-1$ *writes*
- To estimate running time, in classic algorithm theory (ref. 203), we only count the number of operations
- However, if the read and write time is a hundred times the calculation time, then it must be considered
- Without cache, it needs cycle time:
 - $N-1 + 100 * 3*(N-1)$
- With cache and have all data put into it
 - $N-1 + 0.5 * 3*(N-1)$
- When $N=10,000$, the latter is 0.0083 (<1%) of the former!

Cache Policy



- A **Policy** needs to be defined to keep the access to Memory in the cache as much as possible
- $T = m \times T_m + T_h$
- m : cache miss rate
- T_m : memory access time
- T_h : cache access time
- → the lower m the better

First In First Out (FIFO)

Data Addr.

A	B	B	E	F	B	C	B	E
A	A	A	E	E	E	E	E	E
B	B	B	B	F	F	F	F	F
C	C	C	C	C	B	B	B	B
D	D	D	D	D	D	C	C	C

- Missing rate = 4/9
- Always replace the earliest entry in the cache
- This part will be Lab 1

Least recently used (LRU)

Data Addr.

A	B	B	E	F	B	C	B	E
A(0)	A	A	A	A	A	C(7)	C	C
B(-3)	B(2)	B(3)	B	B	B(6)	B	B(8)	B
C(-2)	C	C	E(4)	E	E	E	E	E(9)
D(-1)	D	D	D	F(5)	F	F	F	F

- Missing rate = 3/9
- This strategy requires keeping a "date of birth" for cache lines and tracking the "least recently used" (oldest) cache line based on it
- The difference between LRU and FIFO is that Processor's access to a record will update its "date of birth", while FIFO does not

Random Replacement Policy

- Steps:
 - If the cache is not full, save the most recently accessed data in the cache
 - If it is full, replace one of them at random.
- Simple and stable, once used in ARM processors



Other Policies


- LIFO
- More policies can be found:
- https://en.wikipedia.org/wiki/Cache_replacement_policies

Further Reading

- Cache is everywhere
 - Web Content
 - Video Broadcast
 - Logistic/warehousing
- Can it be solved by ML? Yes.



Amazon Logistics



Policy in Reinforcement Learning

EXPERIMENT

Cache Hit Ratio Optimization



Stefano Tempesta • March 13, 2018



Be the first to like.



Summary

Optimize cache hit ratios and reduce "miss rates" with regression algorithms processed by machine learning.

Performance optimization

- Choose the right algorithm and data structure
 - Go to Course 203
- Choose the right language (C is usually the fastest, but also hard to maintain)
 - Go to Course 102
- Choose compiler
 - icc vs gcc
 - gcc compiler option –O3
 - For python, switch to Cython
- Loop optimization
 - Code movement: extract repeated calculations out of the loop

Performance optimization (serial)

```
1  /* Move call to vec_length out of loop */
2  void combine2(vec_ptr v, data_t *dest)
3  {
4      long int i;
5      long int length = vec_length(v);
6
7      *dest = IDENT;
8      for (i = 0; i < length; i++) {
9          data_t val;
10         get_vec_element(v, i, &val);
11         *dest = *dest OP val;
12     }
13 }
```

- Pay attention to the implicit loop: the linked list must traverse itself when finding the length

Performance optimization

- Write temporary variables inside the loop instead of global variables - compiler can put temporary variables completely in registers without writing back to memory

Row-major order is used in C/C++/Objective-C (for C-style arrays), PL/I,^[4] Pascal,^[5] Speakeasy,^[citation needed] and SAS.^[6]

Column-major order is used in Fortran, MATLAB,^[7] GNU Octave, Julia,^[8] S, S-PLUS,^[9] R,^[10] Scilab,^[11] Yorick, and Rasdaman.^[12]

```
void suma(int rows, int cols, long data[][cols], long *ptotal) {
    long total = 0;
    for (int i = 0; i < rows; i++)
        for (int j = 0; j < cols; j++)
            total += data[i][j];
    *ptotal = total;
}
```

```
void sumb(int rows, int cols, long data[][cols], long *ptotal) {
    long total = 0;
    for (int j = 0; j < cols; j++)
        for (int i = 0; i < rows; i++)
            total += data[i][j];
    *ptotal = total;
}
```

Row-major order

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$

Column-major order

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$

Performance optimization (serial)

```
a@a:~
```

Performance counter stats for './row 10 a' (100 runs):

26.80 msec	task-clock	# 0.988 CPUs utilized	(+- 0.10%)
0	context-switches	# 0.000 /sec	
0	cpu-migrations	# 0.000 /sec	
7,135	page-faults	# 266.118 K/sec	(+- 0.00%)
73,568,049	cycles	# 2.744 GHz	(+- 0.08%) (40.30%)
114,269,204	instructions	# 1.56 insn per cycle	(+- 0.06%) (55.22%)
16,434,432	branches	# 612.963 M/sec	(+- 0.80%) (55.23%)
26,064	branch-misses	# 0.16% of all branches	(+- 0.75%) (55.30%)
8,221,496	L1-dcache-loads	# 306.641 M/sec	(+- 0.07%) (66.42%)
790,271	L1-dcache-load-misses	# 9.59% of all L1-dcache accesses	(+- 1.48%) (74.63%)
39,207	LLC-loads	# 1.462 M/sec	(+- 1.81%) (59.69%)
42,949	LLC-load-misses	# 100.21% of all LLC-cache accesses	(+- 0.80%) (48.44%)

0.0271274 +- 0.0000281 seconds time elapsed (+- 0.10%)

```
a@a:~$ perf stat -d --repeat 100 ./col 10 a
```

Performance counter stats for './col 10 a' (100 runs):

90.95 msec	task-clock	# 1.017 CPUs utilized	(+- 0.20%)
2	context-switches	# 22.459 /sec	(+- 6.24%)
0	cpu-migrations	# 0.000 /sec	
7,134	page-faults	# 80.112 K/sec	(+- 0.00%)
235,965,836	cycles	# 2.650 GHz	(+- 0.07%) (47.24%)
86,014,658	instructions	# 0.37 insn per cycle	(+- 0.11%) (60.41%)
13,131,749	branches	# 147.464 M/sec	(+- 0.26%) (60.42%)
14,802	branch-misses	# 0.12% of all branches	(+- 1.13%) (61.28%)
14,045,201	L1-dcache-loads	# 157.721 M/sec	(+- 0.37%) (65.69%)
8,019,803	L1-dcache-load-misses	# 59.79% of all L1-dcache accesses	(+- 0.14%) (65.94%)
3,652,627	LLC-loads	# 41.017 M/sec	(+- 1.08%) (51.92%)
1,280,277	LLC-load-misses	# 36.37% of all LLC-cache accesses	(+- 1.53%) (47.51%)

0.089411 +- 0.000181 seconds time elapsed (+- 0.20%)

```
a@a:~$
```

CPU has a prefetch mechanism based on locality!


Performance optimization

- Function call
 - Minimize the function call in the loop, or replace it with #define in C or C++ template
- Specify register
 - C++ register variables
- Data type
 - Reduce floating point operations, replace with integers

```
C01E 8D F0    INHEX  BSR    INCH   GET A CHAR
C020 81 30    INHEX  CMP A #'0  ZERO
C022 2B 11    BMI   HEXERR NOT HEX
C024 81 39    CMP A #'9  NINE
C026 2F 0A    BLE   HEXRTS GOOD HEX
C028 81 41    CMP A #'A
C02A 2B 09    BMI   HEXERR NOT HEX
C02C 81 46    CMP A #'F
C02E 2E 05    BGT   HEXERR
C030 80 07    SUB A #7  FIX A-F
C032 84 0F    HEXRTS AND A #$0F CONVERT ASCII TO DIGIT
C034 39      RTS
C035 7E C0 AF  HEXERR JMP   CTRL  RETURN TO CONTROL LOOP
```

The final choice:
Switch to Assembly, if you can!

Further Reading




application-specific integrated circuit

- What is the fastest way to speed up an algorithm?
 - Abandoning the von Neumann architecture
 - Sacrifice Generality for Specialized Acceleration: from CPU to ASIC
 - Example: summing 10,000-length array
 - A circuit supporting 10,000 inputs + 1 output



Field-programmable gate array

- Intelligence Processing Unit @ Graphcore



Design for Deep Learning:
Non-von Neumann architecture
125 Tflops@120W
45TB/s Memory Bandwidth